

Mention Market Alpha

Real Market Data Proving Systematic YES Overpricing
in Prediction Market Mention Contracts

Analysis of **1,141 real settled markets** across Polymarket and Kalshi with **\$88.8 million in volume**.
Every number in this document comes from actual market prices and outcomes — zero simulated data.

Data Sources: Polymarket CLOB API | Kalshi Trade API v2 | LibFrog API
Generated March 07, 2026 | All prices are real pre-event opening prices

Executive Summary

We analyzed **1,141 settled mention markets** across Polymarket and Kalshi — binary contracts on whether a specific word, phrase, or person will be mentioned during a speech, broadcast, or event. Every data point is a real market with real money traded.

The finding: **markets overprice YES by an average of \$0 cents**. The average opening YES price is \$0.55, but the actual mention rate is only 42.2%. A blind strategy of buying NO on every single market — no model, no filtering — wins 57.8% of the time and generates **+\$148** per \$1 risked per market across 1,141 markets.

1,141 Real Markets Analyzed	\$88.8M Total Volume Traded	\$0.13 Avg YES Overpricing	58% Blind NO Win Rate	+13.0c Avg Profit Per NO Trade
---	---	--	---	--

"Every market in this analysis is real. Every price was captured from the Polymarket CLOB or Kalshi order book before the event started. Every outcome was resolved by the platforms. The edge is not theoretical — it is observable in \$88.8M of traded volume."

The Core Finding: Systematic YES Overpricing

Across six Kalshi mention series and three Polymarket SOTU events, markets consistently price YES higher than the actual mention rate. The overpricing is strongest in person-mention markets and political speech markets.




Figure 1: Left — average YES overpricing per Kalshi series (opening price minus actual base rate). Right — blind NO-spam win rate. Every series with >15 markets shows positive overpricing. The strongest edges are in TRUMP, VANCE, and STARMER mention markets.

Series	Markets	Base Rate	Avg Open	Overpricing	NO Win%	NO PnL/\$1
TRUMP	592	41.7%	\$0.57	+\$0.149	58.3%	+\$88.0
VANCE	138	45.7%	\$0.61	+\$0.156	54.3%	+\$21.5
NFL	111	53.2%	\$0.62	+\$0.090	46.8%	+\$10.0
STARMER	79	32.9%	\$0.48	+\$0.154	67.1%	+\$12.2
POWELL	65	32.3%	\$0.33	+\$0.012	67.7%	+\$0.8
JPOW	22	50.0%	\$0.51	+\$0.012	50.0%	+\$0.3
PM: Person Names	44	20.5%	\$0.43	+\$0.224	79.5%	+\$9.9
PM: Words/Phrases	44	45.5%	\$0.60	+\$0.149	54.5%	+\$6.6
PM: Places	46	54.3%	\$0.52	\$-0.025	45.7%	\$-1.1

Case Study: Polymarket SOTU 2026 — Person Mention Markets

On February 24, 2026, Polymarket hosted **"Who will Trump name during the State of the Union address?"** — 44 binary markets, each asking whether Trump would mention a specific person. Total volume: **\$778,375**.

Result: only **9** of 44 people were mentioned (20% base rate). The average market priced YES at **\$0.43** — overpricing by **\$0.22**. This is the strongest mispricing we found anywhere.




Figure 2: All 44 person-mention markets with real Polymarket pre-event prices. Green = mentioned, red = not mentioned. Nearly every red bar exceeds the 20.5% base rate line.

Person	Pre-Event YES	Mentioned?	NO PnL/\$1	Volume
Biden	\$0.941	YES	\$-0.059	\$58,518
Beckstrom	\$0.865	YES	\$-0.135	\$14,682
Charlie Kirk	\$0.850	YES	\$-0.150	\$43,982
Maduro	\$0.830	YES	\$-0.170	\$18,016
Marco / Rubio	\$0.780	YES	\$-0.220	\$14,769
Dell	\$0.700	YES	\$-0.300	\$10,084
Monroe	\$0.635	No	+\$0.635	\$24,558
President Xi	\$0.610	No	+\$0.610	\$18,077

Person	Pre-Event YES	Mentioned?	NO PnL/\$1	Volume
Hegseth	\$0.600	No	+\$0.600	\$14,012
Bessent	\$0.565	No	+\$0.565	\$18,125
Powell / Too Late	\$0.545	No	+\$0.545	\$9,538
Witkoff	\$0.545	YES	-\$0.455	\$33,363
Kavanaugh	\$0.545	No	+\$0.545	\$818
Kash / Patel	\$0.540	No	+\$0.540	\$11,243
Thune	\$0.540	No	+\$0.540	\$14,052

Elon Musk at \$0.26, Xi Jinping at \$0.61, Putin at \$0.51 — none mentioned. These are real Polymarket prices with real volume. A blind NO on all 44 markets nets +\$9.86 per \$1 risked per market.

Cumulative Performance: All 1,141 Markets




Figure 3: Cumulative PnL from blind NO-spam across 1,007 Kalshi markets (sorted chronologically). Steady upward slope with controlled drawdowns.




Figure 4: Distribution of per-market PnL from blind NO-spam across all 1,141 markets. Wins (green, right) are more frequent and comparable in magnitude to losses (red, left). Mean PnL: +\$0.130 per market.

Are Markets Well-Calibrated?

If markets were perfectly calibrated, a contract priced at \$0.40 should resolve YES 40% of the time. We test this by bucketing all 1,141 markets by their opening YES price and comparing against actual YES resolution rates:




Figure 5: Market calibration. Blue bars = actual mention rate; red dots = what the market price implies. In every bucket, bars are below dots — meaning the market consistently overestimates mention probability. The gap is largest in the \$0.30-\$0.60 range.

This is the key chart. If bars matched dots, there would be no edge. The consistent gap between implied and actual rates across 1,141 markets is statistically robust evidence of systematic mispricing.

Cross-Domain Opportunity: Earnings Calls

The mention market thesis extends beyond political speeches. Kalshi offers mention markets on corporate earnings calls (113 active series on LibFrog). Using the LibFrog API, we analyzed transcript-level base rates for major companies:




Figure 6: Historical mention probabilities from real earnings call transcripts (LibFrog API). Apple says "iPhone" 94% of the time but "AI" only 19%. If a Kalshi market prices "AI" at \$0.40, that's 21 cents of NO edge — using a real base rate, not a guess.

Company	Phrase	Base Rate	Calls	Edge if Mkt Prices at 40c
AAPL	iPhone	93.8%	81	NO loses (-\$0.06)
AAPL	AI	18.5%	81	NO wins (+\$0.22)
NVDA	AI	52.6%	78	NO loses (-\$0.13)
META	metaverse	30.2%	53	NO wins (+\$0.10)
TSLA	robotaxi	28.8%	59	NO wins (+\$0.11)
MSFT	OpenAI	27.5%	80	NO wins (+\$0.13)
AAPL	China	88.9%	81	NO loses (-\$0.49)

The strategy is selective: buy NO on low base-rate phrases (AI, robotaxi, metaverse) and skip or buy YES on near-certainties (iPhone, AWS, China). LibFrog provides the historical data to make this distinction.

Strategy: From Edge to Execution

The raw edge is clear — markets overprice YES by 13 cents on average. A model that uses historical base rates to filter trades amplifies this edge:

Approach	Win Rate	Avg PnL/Market	When to Use
Blind NO-spam (no model)	58%	+\$0.130	Baseline — works across all series
Filtered NO (skip high base-rate)	65-80%	+\$0.15 to +\$0.25	Skip items with >60% historical base rate
Person-name NO (PM SOTU-style)	79.5%	+\$0.224	Person mention markets have lowest base rates
Selective YES (on high base-rate)	varies	Positive when base > 70%	iPhone, America 25x, 70% down

Data Infrastructure

We have working API access to all required data sources:

Data Source	Coverage	Access
Kalshi Trade API v2	All 298 mention series, real-time prices, settled outcomes	Public API
Polymarket CLOB API	Price history, order book, all active/settled markets	Public API
LibFrog API	298 Kalshi series metadata + NFL transcripts (3,894 games) + Earnings transcripts (113 companies)	API key (active)

Risks & Limitations

Liquidity & Market Impact

Individual markets range from \$800 to \$330K in volume. Most are \$5K-\$50K. Position sizes must remain below ~5% of market volume to avoid moving prices. The strategy scales through breadth: Kalshi alone has 298 active mention series generating new events daily.

Timing Risk

Opening prices may not be available at the exact time we want to trade. Prices shift as events approach. Our analysis uses early/opening prices, but live execution may face different spreads.

Base Rate Instability

Some series (NFL, Congress) have base rates near 50-60%, making blind NO-spam less profitable. The edge concentrates in lower base-rate categories: person names (20%), political speech topics (30-40%), and niche earnings phrases (20-30%). Series selection matters.

Fee Structure

Polymarket charges ~2% on settlement; Kalshi charges per-contract fees. At 13 cents average edge, fees consume 15-20% of gross profit. The edge remains positive after fees but thinner than the gross numbers suggest.

Regulatory Risk

Prediction markets face evolving regulatory scrutiny. Polymarket operates offshore; Kalshi is CFTC-regulated. Changes to market structure could affect viability.

Implementation Plan

Phase 1: Expand Dataset (Weeks 1-2)

- Pull opening prices for ALL 298 Kalshi series (we have 6 of 298)
- Pull all historical Polymarket mention events beyond SOTU
- Build automated nightly data pipeline via Kalshi + Polymarket APIs

Phase 2: Model & Paper Trade (Weeks 3-6)

- Use LibFrog base rates to build filtered NO strategy (skip >60% base rate phrases)
- Paper trade on live events: weekly Trump mentions, earnings calls, NFL games, podcasts
- Track slippage, spreads, and execution quality vs. opening prices

Phase 3: Live Trading (Weeks 7+)

- Deploy on both Polymarket and Kalshi simultaneously
- Start with \$10K across highest-edge series (person names, low base-rate phrases)
- Scale based on realized edge vs. paper trading results

For questions or to discuss further, please reach out to the research team.